

From-scratch Network

May 2, 2022

1 Deep Learning From-scratch Network Training

Zach Kangas

5/2/22

Dr. Yoder

Deep Learning

2 Introduction

In this report, we look at using the from-scratch neural network created with PyTorch to train on a linear toy dataset, then the CIFAR-10 and MNIST flattened datasets. The results from the training of each dataset is described after training, in a final validation statistic and a graph of the loss and accuracy of the training and validation sets.

2.1 Importing Relevant Libraries

```
[1]: import torch
import numpy as np
import matplotlib.pyplot as plt
import torchvision
import warnings
import os.path
import math
import seaborn as sns

import NeuralNetwork
import Client
```

2.2 Setting up dtype, device, and datapaths

```
[2]: # warnings.filterwarnings('ignore') # If you see warnings that you know you
    ↪ can ignore, it can be useful to enable this.

# For fashion-MNIST and similar problems
```

```

DATA_ROOT = '/data/cs3450/data/'
FASHION_MNIST_TRAINING = '/data/cs3450/data/fashion_mnist_flattened_training.
    ↪npz'
FASHION_MNIST_TESTING = '/data/cs3450/data/fashion_mnist_flattened_testing.npz'
CIFAR10_TRAINING = '/data/cs3450/data/cifar10_flattened_training.npz'
CIFAR10_TESTING = '/data/cs3450/data/cifar10_flattened_testing.npz'
CIFAR100_TRAINING = '/data/cs3450/data/cifar100_flattened_training.npz'
CIFAR100_TESTING = '/data/cs3450/data/cifar100_flattened_testing.npz'

DTYPE = torch.float32
# With this block, we don't need to set device=DEVICE for every tensor.
torch.set_default_dtype(torch.float32)
if torch.cuda.is_available():
    torch.cuda.set_device(0)
    torch.set_default_tensor_type(torch.cuda.FloatTensor)
    print("Running on the GPU")
else:
    print("Running on the CPU")

```

Running on the GPU

2.3 Importing training data

```

[3]: def create_linear_training_data(training_points):
    """
    This method simply rotates points in a 2D space.
    Be sure to use L2 regression in the place of the final softmax layer before
    ↪testing on this
    data!
    :return: (x,y) the dataset. x is a numpy array where columns are training
    ↪samples and
            y is a numpy array where columns are one-hot labels for the
    ↪training sample.
    """
    x = torch.randn((2, training_points))
    x1 = x[0:1, :].clone()
    x2 = x[1:2, :]
    y = torch.cat((-x2, x1), axis=0)
    return x, y

def create_folded_training_data():
    """
    This method introduces a single non-linear fold into the sort of data
    ↪created by create_linear_training_data. Be sure to REMOVE the final softmax
    ↪layer before testing on this data!
    Be sure to use L2 regression in the place of the final softmax layer before
    ↪testing on this
    """

```

```

    data!
    :return: (x,y) the dataset. x is a numpy array where columns are training_
    ↪samples and
           y is a numpy array where columns are one-hot labels for the_
    ↪training sample.
    """
    x = torch.randn((2, TRAINING_POINTS))
    x1 = x[0:1, :].clone()
    x2 = x[1:2, :]
    x2 *= 2 * ((x2 > 0).float() - 0.5)
    y = torch.cat((-x2, x1), axis=0)
    return x, y

def create_square():
    """
    This is a square example
    insideness is true if the points are inside the square.
    :return: (points, insideness) the dataset. points is a 2xN array of points_
    ↪and insideness is true if the point is inside the square.
    """
    win_x = [2,2,3,3]
    win_y = [1,2,2,1]
    win = torch.tensor([win_x,win_y],dtype=torch.float32)
    win_rot = torch.cat((win[:,1:],win[:,0:1]),axis=1)
    t = win_rot - win # edges tangent along side of poly
    rotation = torch.tensor([[0, 1],[-1,0]],dtype=torch.float32)
    normal = rotation @ t # normal vectors to each side of poly
    # torch.matmul(rotation,t) # Same thing

    points = torch.rand((2,2000),dtype = torch.float32)
    points = 4*points

    vectors = points[:,np.newaxis,:] - win[:, :,np.newaxis] # reshape to fill_
    ↪origin
    insideness = (normal[:, :,np.newaxis] * vectors).sum(axis=0)
    insideness = insideness.T
    insideness = insideness > 0
    insideness = insideness.all(axis=1)
    return points, insideness

def create_patterns():
    """
    I don't remember what sort of data this generates -- Dr. Yoder

```

```

        :return: (points, insideness) the dataset. points is a 2xN array of points,
        ↪ and insideness is true if the point is inside the square.
        """
        pattern1 = torch.tensor([[1, 0, 1, 0, 1, 0]], dtype=torch.float32).T
        pattern2 = torch.tensor([[1, 1, 1, 0, 0, 0]], dtype=torch.float32).T
        num_samples = 1000

        x = torch.zeros((pattern1.shape[0], num_samples))
        y = torch.zeros((2, num_samples))
        # TODO: Implement with shuffling instead?
        for i in range(0, num_samples):
            if torch.rand(1) > 0.5:
                x[:, i:i+1] = pattern1
                y[:, i:i+1] = torch.tensor([[0, 1]], dtype=torch.float32).T
            else:
                x[:, i:i+1] = pattern2
                y[:, i:i+1] = torch.tensor([[1, 0]], dtype=torch.float32).T
        return x, y

def load_dataset_flattened(train=True, dataset='Fashion-MNIST', download=False):
    """
    :param train: True for training, False for testing
    :param dataset: 'Fashion-MNIST', 'CIFAR-10', or 'CIFAR-100'
    :param download: True to download. Keep to false afterwards to avoid
    ↪ unneeded downloads.
    :return: (x,y) the dataset. x is a numpy array where columns are training
    ↪ samples and
           y is a numpy array where columns are one-hot labels for the
    ↪ training sample.
    """
    if dataset == 'Fashion-MNIST':
        if train:
            path = FASHION_MNIST_TRAINING
        else:
            path = FASHION_MNIST_TESTING
        num_labels = 10
    elif dataset == 'CIFAR-10':
        if train:
            path = CIFAR10_TRAINING
        else:
            path = CIFAR10_TESTING
        num_labels = 10
    elif dataset == 'CIFAR-100':
        if train:
            path = CIFAR100_TRAINING
        else:

```

```

        path = CIFAR100_TESTING
        num_labels = 100
    else:
        raise ValueError('Unknown dataset: '+str(dataset))

    if os.path.isfile(path):
        print('Loading cached flattened data for',dataset,'training' if train_
→else 'testing')
        data = np.load(path)
        x = torch.tensor(data['x'],dtype=torch.float32)
        y = torch.tensor(data['y'],dtype=torch.float32)
        pass
    else:
        class ToTorch(object):
            """Like ToTensor, only to a numpy array"""

            def __call__(self, pic):
                return torchvision.transforms.functional.to_tensor(pic)

        if dataset == 'Fashion-MNIST':
            data = torchvision.datasets.FashionMNIST(
                root=DATA_ROOT, train=train, transform=ToTorch(),
→download=download)
        elif dataset == 'CIFAR-10':
            data = torchvision.datasets.CIFAR10(
                root=DATA_ROOT, train=train, transform=ToTorch(),
→download=download)
        elif dataset == 'CIFAR-100':
            data = torchvision.datasets.CIFAR100(
                root=DATA_ROOT, train=train, transform=ToTorch(),
→download=download)
        else:
            raise ValueError('This code should be unreachable because of a
→previous check.')
        x = torch.zeros((len(data[0][0].flatten()), len(data)),dtype=torch.
→float32)
        for index, image in enumerate(data):
            x[:, index] = data[index][0].flatten()
        labels = torch.tensor([sample[1] for sample in data])
        y = torch.zeros((num_labels, len(labels)), dtype=torch.float32)
        y[labels, torch.arange(len(labels))] = 1
        np.savez(path, x=x.detach().numpy(), y=y.detach().numpy())
    return x, y

```

3 Training the Network - Linear

In this section, I use a small network to predict on the linear dataset. This really should not need any more than 3-5 nodes in the hidden layer, but I chose 10 for this example as the train time was extremely quick regardless.

```
[4]: x, y = create_linear_training_data(10000)
epochs = 10

nn = NeuralNetwork.Network(x.shape[0], y.shape[0], dtype=torch.float32,
    ↳loss="l2", regularization_factor=0.01, learning_rate=0.001)
lin1 = nn.add_linear_generated(num_nodes=10, w=0.05, wo=0, b=0, bo=0,
    ↳regularization=True)
rel1 = nn.add_relu()
lin2 = nn.add_linear_generated(num_nodes=2, w=0.05, wo=0, b=0, bo=0,
    ↳regularization=True)

client = Client.Client(nn, x, y)
tloss, tacc = client.train_stats()
vloss, vacc = client.validation_stats()
print("E: -1", "\ttL:", tloss, "\ttA:", tacc, "\tvL:", vloss, "\tvA:", vacc)
train_data = client.train(epochs, 1, verbose=True)
```

```
E: -1    tL: 1.9948890209197998  tA: 0.4767500162124634  vL: 1.882861614227295
vA: 0.5045000314712524
E: 0     tL: 0.0788913443684578  tA: 0.9827500581741333  vL: 0.06811783462762833
vA: 0.9775000214576721
E: 1     tL: 0.02252291329205036          tA: 0.9888750314712524  vL:
0.019315805286169052          vA: 0.9900000691413879
E: 2     tL: 0.013064330443739891          tA: 0.9937500357627869  vL:
0.010838055051863194          vA: 0.9915000200271606
E: 3     tL: 0.008565248921513557          tA: 0.9961250424385071  vL:
0.007008089683949947          vA: 0.9965000748634338
E: 4     tL: 0.006365333218127489          tA: 0.9978750348091125  vL:
0.005219032522290945          vA: 0.9985000491142273
E: 5     tL: 0.005125655326992273          tA: 0.9982500672340393  vL:
0.0042310841381549835          vA: 0.9985000491142273
E: 6     tL: 0.004361961502581835          tA: 0.9986250400543213  vL:
0.0036184692289680243          vA: 0.9985000491142273
E: 7     tL: 0.003849930362775922          tA: 0.9987500309944153  vL:
0.0031996257603168488          vA: 0.9985000491142273
E: 8     tL: 0.003465912537649274          tA: 0.9987500309944153  vL:
0.0028836484998464584          vA: 0.9985000491142273
E: 9     tL: 0.0031606603879481554          tA: 0.9988750219345093  vL:
0.0026310631074011326          vA: 0.9985000491142273
```

3.1 Statistics

Final Training Loss*: 0.003

Final Training Accuracy*: 0.999


Final Validation Loss: 0.003

Final Validation Accuracy: 0.999

*: Validation and training are basically identical

3.1.1 Training Graph

```
[5]: fig = plt.figure(figsize=(10,10))
pls = np.array(train_data)
ax = sns.lineplot(x=pls[:,0], y=pls[:,1], color="lightcoral", label="Train Loss")
ax = sns.lineplot(x=pls[:,0], y=pls[:,3], color="r", label="Validation Loss")
ax.set_xlabel("Epochs")
ax.set_ylabel("Loss")
ax2 = plt.twinx()
ax2.set_ylabel("Accuracy")
sns.lineplot(x=pls[:,0], y=pls[:,2], color="lightsteelblue", label="Train Acc")
sns.lineplot(x=pls[:,0], y=pls[:,4], color="b", label="Validation Acc")
plt.show()
```



4 Training the Network - MNIST

```
[6]: dataset = 'Fashion-MNIST'

x, y = load_dataset_flattened(train=True, dataset=dataset, download=False)
epochs = 50

nn = NeuralNetwork.Network(x.shape[0], y.shape[0], loss="ce",
    ↳regularization_factor=0.00001, learning_rate=0.001)
lin3 = nn.add_linear_generated(num_nodes=800, w=0.001, wo=0, b=0, bo=0,
    ↳regularization=True)
rel = nn.add_relu()
lin3 = nn.add_linear_generated(num_nodes=300, w=0.0025, wo=0, b=0, bo=0,
    ↳regularization=True)
rel = nn.add_relu()
```



```

lin4 = nn.add_linear_generated(num_nodes=y.shape[0], w=0.03, wo=0, b=0, bo=0,
    ↪regularization=True)
nn.add_softmax()

client = Client.Client(nn, x, y)
tloss, tacc = client.train_stats()
vloss, vacc = client.validation_stats()
print("E: -1", "\ttL:", tloss, "\ttA:", tacc, "\tvL:", vloss, "\tvA:", vacc)
train_data = client.train(epochs, 1, verbose=True)

```

Loading cached flattened data for Fashion-MNIST training

```

E: -1   tL: 2.301969289779663   tA: 0.10106249898672104   vL:
2.3022379875183105   vA: 0.09574999660253525
E: 0    tL: 1.0199207067489624   tA: 0.6098541617393494   vL: 1.0165126323699951
vA: 0.6106666326522827
E: 1    tL: 0.7970491647720337   tA: 0.7301666736602783   vL: 0.8024205565452576
vA: 0.7242500185966492
E: 2    tL: 0.6782697439193726   tA: 0.7927708029747009   vL: 0.6892938613891602
vA: 0.7858332991600037
E: 3    tL: 0.5547126531600952   tA: 0.835979163646698   vL: 0.5694445371627808
vA: 0.828249990940094
E: 4    tL: 0.5016655921936035   tA: 0.8491041660308838   vL: 0.520009458065033
vA: 0.8395833373069763
E: 5    tL: 0.4686982035636902   tA: 0.8564791679382324   vL: 0.49052560329437256
vA: 0.8460000157356262
E: 6    tL: 0.4433543384075165   tA: 0.8630833029747009   vL: 0.46792879700660706
vA: 0.8528333306312561
E: 7    tL: 0.42368823289871216   tA: 0.867354154586792   vL:
0.4514448046684265   vA: 0.856249988079071
E: 8    tL: 0.4073924124240875   tA: 0.871749997138977   vL: 0.43798235058784485
vA: 0.859749972820282
E: 9    tL: 0.3929104804992676   tA: 0.8749791383743286   vL: 0.4260786175727844
vA: 0.8619999885559082
E: 10   tL: 0.3799373209476471   tA: 0.8788958191871643   vL: 0.41609740257263184
vA: 0.8644166588783264
E: 11   tL: 0.3679186701774597   tA: 0.883104145526886   vL: 0.40698710083961487
vA: 0.8667500019073486
E: 12   tL: 0.3589191138744354   tA: 0.8858749866485596   vL: 0.40038353204727173
vA: 0.8667500019073486
E: 13   tL: 0.3506874144077301   tA: 0.8873957991600037   vL: 0.3951770067214966
vA: 0.8680832982063293
E: 14   tL: 0.3408052623271942   tA: 0.890625   vL: 0.38824355602264404
vA: 0.8704166412353516
E: 15   tL: 0.33359503746032715   tA: 0.8930000066757202   vL:
0.3835214078426361   vA: 0.8709999918937683
E: 16   tL: 0.32471516728401184   tA: 0.895145833492279   vL:
0.3775201737880707   vA: 0.8725833296775818

```

E: 17 tL: 0.3182857632637024 tA: 0.8966875076293945 vL: 0.3744330108165741
 vA: 0.8734166622161865
 E: 18 tL: 0.3099108934402466 tA: 0.8997499942779541 vL: 0.3686022162437439
 vA: 0.8758333325386047
 E: 19 tL: 0.30336180329322815 tA: 0.9010416865348816 vL:
 0.36507558822631836 vA: 0.8761666417121887
 E: 20 tL: 0.29601964354515076 tA: 0.9032708406448364 vL:
 0.360989511013031 vA: 0.8775833249092102
 E: 21 tL: 0.2914585769176483 tA: 0.9045208096504211 vL: 0.3592927157878876
 vA: 0.878083348274231
 E: 22 tL: 0.2851117253303528 tA: 0.9055416584014893 vL: 0.3562021553516388
 vA: 0.8774999976158142
 E: 23 tL: 0.2775116264820099 tA: 0.909333348274231 vL: 0.35242822766304016
 vA: 0.8779166340827942
 E: 24 tL: 0.27286866307258606 tA: 0.9104791283607483 vL:
 0.3509041368961334 vA: 0.8787499666213989
 E: 25 tL: 0.2669496536254883 tA: 0.9120833277702332 vL: 0.3480924963951111
 vA: 0.8808333277702332
 E: 26 tL: 0.2623220384120941 tA: 0.9133749604225159 vL: 0.3467816114425659
 vA: 0.8814166784286499
 E: 27 tL: 0.2570047080516815 tA: 0.914354145526886 vL: 0.3445451259613037
 vA: 0.8798333406448364
 E: 28 tL: 0.25267696380615234 tA: 0.9165208339691162 vL:
 0.3447311222553253 vA: 0.8799166679382324
 E: 29 tL: 0.24646137654781342 tA: 0.9169583320617676 vL:
 0.3419124186038971 vA: 0.8807500004768372
 E: 30 tL: 0.24111855030059814 tA: 0.9186874628067017 vL:
 0.33942484855651855 vA: 0.8819166421890259
 E: 31 tL: 0.24070051312446594 tA: 0.918999969959259 vL:
 0.3421674072742462 vA: 0.8805000185966492
 E: 32 tL: 0.23306208848953247 tA: 0.9228541851043701 vL:
 0.337630033493042 vA: 0.8828333020210266
 E: 33 tL: 0.2310403287410736 tA: 0.9213333129882812 vL: 0.3385462164878845
 vA: 0.8816666603088379
 E: 34 tL: 0.22413130104541779 tA: 0.9251874685287476 vL:
 0.3358246684074402 vA: 0.8844999670982361
 E: 35 tL: 0.22544534504413605 tA: 0.9226458072662354 vL:
 0.3413766622543335 vA: 0.8818333148956299
 E: 36 tL: 0.21809890866279602 tA: 0.9243957996368408 vL:
 0.33726605772972107 vA: 0.8799999952316284
 E: 37 tL: 0.21374984085559845 tA: 0.9277708530426025 vL:
 0.33766958117485046 vA: 0.8824166655540466
 E: 38 tL: 0.2118811011314392 tA: 0.9261249899864197 vL: 0.34009861946105957
 vA: 0.8815833330154419
 E: 39 tL: 0.2011374831199646 tA: 0.929562509059906 vL: 0.3334323763847351
 vA: 0.8803333044052124
 E: 40 tL: 0.19780659675598145 tA: 0.9322291612625122 vL:
 0.33363547921180725 vA: 0.8814166784286499

E: 41 tL: 0.19463203847408295 tA: 0.9322708249092102 vL:
 0.33499154448509216 vA: 0.8822500109672546
 E: 42 tL: 0.1957651823759079 tA: 0.9289374947547913 vL: 0.3398110568523407
 vA: 0.8805833458900452
 E: 43 tL: 0.2429502159357071 tA: 0.9192500114440918 vL: 0.34988078474998474
 vA: 0.8775833249092102
 E: 44 tL: 0.21054844558238983 tA: 0.9306041598320007 vL:
 0.3362560272216797 vA: 0.8819999694824219
 E: 45 tL: 0.19667980074882507 tA: 0.9354791641235352 vL:
 0.33261504769325256 vA: 0.8823333382606506
 E: 46 tL: 0.19415238499641418 tA: 0.9334166646003723 vL:
 0.3360462486743927 vA: 0.8822500109672546
 E: 47 tL: 0.1902884840965271 tA: 0.9346041679382324 vL: 0.33751821517944336
 vA: 0.8814166784286499
 E: 48 tL: 0.18186955153942108 tA: 0.937874972820282 vL:
 0.3320619761943817 vA: 0.8849999904632568
 E: 49 tL: 0.18418346345424652 tA: 0.9369791746139526 vL:
 0.3405294418334961 vA: 0.8833333253860474

4.1 Statistics

Final Training Loss: 0.184


Final Training Accuracy: 0.937

Final Validation Loss: 0.341

Final Validation Accuracy: 0.883

4.1.1 Training Graph

```
[7]: fig = plt.figure(figsize=(10,10))
pls = np.array(train_data)
ax = sns.lineplot(x=pls[:,0], y=pls[:,1], color="lightcoral", label="Train Loss")
ax = sns.lineplot(x=pls[:,0], y=pls[:,3], color="r", label="Validation Loss")
ax.set_xlabel("Epochs")
ax.set_ylabel("Loss")
ax2 = plt.twinx()
ax2.set_ylabel("Accuracy")
sns.lineplot(x=pls[:,0], y=pls[:,2], color="lightsteelblue", label="Train Acc")
sns.lineplot(x=pls[:,0], y=pls[:,4], color="b", label="Validation Acc")
plt.show()
```



5 Training the Network - CIFAR-10

```
[8]: %%time
dataset = 'CIFAR-10'

x, y = load_dataset_flattened(train=True, dataset=dataset, download=False)
epochs = 50

nn = NeuralNetwork.Network(x.shape[0], y.shape[0], loss="ce",
    ↳regularization_factor=0.0001, learning_rate=0.0002)
lin3 = nn.add_linear_generated(num_nodes=800, w=0.001, wo=0, b=0, bo=0,
    ↳regularization=True)
rel = nn.add_relu()
lin3 = nn.add_linear_generated(num_nodes=300, w=0.0025, wo=0, b=0, bo=0,
    ↳regularization=True)
```

```

rel = nn.add_relu()
lin4 = nn.add_linear_generated(num_nodes=y.shape[0], w=0.03, wo=0, b=0, bo=0,
    ↪regularization=True)
nn.add_softmax()

client = Client.Client(nn, x, y)

tloss, tacc = client.train_stats()
vloss, vacc = client.validation_stats()
print("E: -1", "\ttL:", tloss, "\ttA:", tacc, "\tvL:", vloss, "\tvA:", vacc)
train_data = client.train(epochs, 1, verbose=True)

```

Loading cached flattened data for CIFAR-10 training

```

E: -1   tL: 2.306262493133545   tA: 0.10119999945163727   vL:
2.306717872619629   vA: 0.09519999474287033
E: 0    tL: 2.278472900390625   tA: 0.14217498898506165   vL:
2.277625799179077   vA: 0.1444000005722046
E: 1    tL: 2.0832886695861816   tA: 0.21222499012947083   vL:
2.0869710445404053   vA: 0.21359999477863312
E: 2    tL: 2.0263068675994873   tA: 0.2530499994754791   vL: 2.0331838130950928
vA: 0.25119999051094055
E: 3    tL: 1.968971848487854   tA: 0.2750999927520752   vL: 1.9787909984588623
vA: 0.2685999870300293
E: 4    tL: 1.9420374631881714   tA: 0.29794999957084656   vL:
1.9531933069229126   vA: 0.2920999825000763
E: 5    tL: 1.9121125936508179   tA: 0.3107749819755554   vL: 1.924224615097046
vA: 0.30410000681877136
E: 6    tL: 1.8649016618728638   tA: 0.32850000262260437   vL:
1.8782538175582886   vA: 0.3212999999523163
E: 7    tL: 1.8164889812469482   tA: 0.3481749892234802   vL: 1.8317759037017822
vA: 0.3409000039100647
E: 8    tL: 1.7772924900054932   tA: 0.3669999837875366   vL: 1.7953062057495117
vA: 0.3572999835014343
E: 9    tL: 1.7397589683532715   tA: 0.3808249831199646   vL: 1.7609741687774658
vA: 0.37209999561309814
E: 10   tL: 1.7067975997924805   tA: 0.3930499851703644   vL: 1.7314820289611816
vA: 0.38449999690055847
E: 11   tL: 1.6775034666061401   tA: 0.40414997935295105   vL:
1.7058568000793457   vA: 0.3905999958515167
E: 12   tL: 1.650678277015686   tA: 0.41349998116493225   vL:
1.6829472780227661   vA: 0.40059998631477356
E: 13   tL: 1.6289818286895752   tA: 0.42114999890327454   vL:
1.6646416187286377   vA: 0.407399982213974
E: 14   tL: 1.6080230474472046   tA: 0.4297249913215637   vL: 1.6473249197006226
vA: 0.41449999809265137
E: 15   tL: 1.5864176750183105   tA: 0.4385499954223633   vL: 1.6296629905700684
vA: 0.42319998145103455

```

E: 16 tL: 1.5719285011291504 tA: 0.44347497820854187 vL:
 1.6187530755996704 vA: 0.42579999566078186
 E: 17 tL: 1.5556085109710693 tA: 0.45159998536109924 vL:
 1.6060779094696045 vA: 0.4324999749660492
 E: 18 tL: 1.5399166345596313 tA: 0.4574749767780304 vL: 1.5948086977005005
 vA: 0.44039997458457947
 E: 19 tL: 1.5249499082565308 tA: 0.46265000104904175 vL:
 1.58340322971344 vA: 0.4438999891281128
 E: 20 tL: 1.5090336799621582 tA: 0.4682749807834625 vL: 1.572184681892395
 vA: 0.44829997420310974
 E: 21 tL: 1.4942880868911743 tA: 0.47269999980926514 vL:
 1.5616765022277832 vA: 0.44849997758865356
 E: 22 tL: 1.4808235168457031 tA: 0.4781249761581421 vL: 1.5525765419006348
 vA: 0.4519999921321869
 E: 23 tL: 1.4645767211914062 tA: 0.4840250015258789 vL: 1.5414317846298218
 vA: 0.45719999074935913
 E: 24 tL: 1.4479894638061523 tA: 0.4894999861717224 vL: 1.5305671691894531
 vA: 0.46149998903274536
 E: 25 tL: 1.4361202716827393 tA: 0.493149995803833 vL: 1.5239161252975464
 vA: 0.4624999761581421
 E: 26 tL: 1.4188899993896484 tA: 0.5005499720573425 vL: 1.5127066373825073
 vA: 0.4657000005245209
 E: 27 tL: 1.404828429222107 tA: 0.5052750110626221 vL: 1.5047402381896973
 vA: 0.47119998931884766
 E: 28 tL: 1.3929346799850464 tA: 0.5085999965667725 vL: 1.4980820417404175
 vA: 0.4721999764442444
 E: 29 tL: 1.3696414232254028 tA: 0.5182749629020691 vL: 1.4816571474075317
 vA: 0.47859999537467957
 E: 30 tL: 1.3558194637298584 tA: 0.5236250162124634 vL: 1.4731965065002441
 vA: 0.48069998621940613
 E: 31 tL: 1.339540958404541 tA: 0.5292750000953674 vL: 1.4652029275894165
 vA: 0.48429998755455017
 E: 32 tL: 1.327163815498352 tA: 0.5330749750137329 vL: 1.4585258960723877
 vA: 0.487199991941452
 E: 33 tL: 1.3164252042770386 tA: 0.5364499688148499 vL: 1.4545962810516357
 vA: 0.48799997568130493
 E: 34 tL: 1.298164963722229 tA: 0.5413749814033508 vL: 1.4455015659332275
 vA: 0.4916999936103821
 E: 35 tL: 1.2883391380310059 tA: 0.5440499782562256 vL: 1.4426584243774414
 vA: 0.4912000000476837
 E: 36 tL: 1.2753905057907104 tA: 0.548799991607666 vL: 1.4392447471618652
 vA: 0.4921000003814697
 E: 37 tL: 1.2615435123443604 tA: 0.5529749989509583 vL: 1.4343303442001343
 vA: 0.4940999746322632
 E: 38 tL: 1.2520612478256226 tA: 0.5554749965667725 vL: 1.4325724840164185
 vA: 0.49149999022483826
 E: 39 tL: 1.2406355142593384 tA: 0.5584749579429626 vL: 1.4311285018920898
 vA: 0.4918999969959259

E: 40 tL: 1.2233253717422485 tA: 0.5666999816894531 vL: 1.4223697185516357
 vA: 0.49609997868537903
 E: 41 tL: 1.2063863277435303 tA: 0.5703749656677246 vL: 1.4170548915863037
 vA: 0.49559998512268066
 E: 42 tL: 1.197527527809143 tA: 0.5738499760627747 vL: 1.4165488481521606
 vA: 0.4982999861240387
 E: 43 tL: 1.190682053565979 tA: 0.5746999979019165 vL: 1.420558214187622
 vA: 0.49449998140335083
 E: 44 tL: 1.1825757026672363 tA: 0.5766249895095825 vL: 1.420655608177185
 vA: 0.49469998478889465
 E: 45 tL: 1.1709179878234863 tA: 0.5788750052452087 vL: 1.4202147722244263
 vA: 0.49459999799728394
 E: 46 tL: 1.1652510166168213 tA: 0.5796499848365784 vL: 1.4253321886062622
 vA: 0.4948999881744385
 E: 47 tL: 1.1514488458633423 tA: 0.5859999656677246 vL: 1.4215196371078491
 vA: 0.49539998173713684
 E: 48 tL: 1.148616075515747 tA: 0.5854249596595764 vL: 1.4280489683151245
 vA: 0.49629998207092285
 E: 49 tL: 1.1397567987442017 tA: 0.5858500003814697 vL: 1.4298659563064575
 vA: 0.4948999881744385
 CPU times: user 36min 29s, sys: 1min 7s, total: 37min 37s
 Wall time: 37min 37s

5.1 Statistics

Final Training Loss: 1.140


Final Training Accuracy: 0.586

Final Validation Loss: 1.43

Final Validation Accuracy: 0.495

5.1.1 Training Graph

```
[9]: fig = plt.figure(figsize=(10,10))
    pls = np.array(train_data)
    ax = sns.lineplot(x=pls[:,0], y=pls[:,1], color="lightcoral", label="Train Loss")
    ax = sns.lineplot(x=pls[:,0], y=pls[:,3], color="r", label="Validation Loss")
    ax.set_xlabel("Epochs")
    ax.set_ylabel("Loss")
    ax2 = plt.twinx()
    ax2.set_ylabel("Accuracy")
    sns.lineplot(x=pls[:,0], y=pls[:,2], color="lightsteelblue", label="Train Acc")
    sns.lineplot(x=pls[:,0], y=pls[:,4], color="b", label="Validation Acc")
    plt.show()
```



6 Reflection / Discussion

6.1 Start of the network

The initial lab felt very odd in this series. The client, how it was suggested to be implemented, still doesn't quite make sense to me. Hence I have likely a different client implementation than most other people.

6.1.1 Design Philosophy

The suggestion to manually create the layers and link them together I think in some ways was a fine one. My lab 3 I believe had basically this feature. But very quickly I struggled with modifying the layers, linking them to the output, and deciding what to do with them. So what I ended up doing was simply making a driver, in the form of the NeuralNetwork class, to handle that for me.

My gradient tape was a linked list with manual calls to the regularization layers. This made more sense to me than doing an array, which I did in lab 3 and was one of the reasons why I thought

that lab to be very tedious.

Because of this, I certainly had obfuscated information. It would be neigh impossible to call a layer mid-inference, and difficult to show the structure of the network without just showing the code that made it.

But because the wrappers were so comprehensive in my network, it meant I can feasibly train on any dataset with any amount of layers, types of layers, and size of layers without much issue. Adding types layers is just adding the layer in layers then adding a wrapper, easy as that.

My client was basically what you'd get with a scikit learn wrapper. Train, train test split, inference, and statistics. Not much more, not much less.

6.2 Back Prop

I thought the experience deriving the back propogation was fine. I do wish there was a little more guidance, and, moreso than guidance, specific suggestions on my original backprop that would say 'this spot is wrong or will be hard to translate due to unclear stuff.' I saw that first hand when I began unit testing, especially against pytorch's autograd, where and how some of my gradients were just minimally wrong.

More time needed to be put into a single layer backprop instead of going through the whole network in week 3-5 and calculating each layer in a sludge of information. My fist quiz on backprop I failed (9/20?) yet the second, which was a test, I think I nearly aced. It was connecting the dots from the sludge of information where providing the dots would have been more effective.

6.3 Unit Tests

The unit tests have me split. Yes, it is good to do them. But at the same time, my unit tests I am turning in now *do not work*. It is not because my equations are wrong, but it's because I had to change a couple things from the toy unit tests to an actual training scenerio in my Network class that broke the unit tests. Could I have seen that coming, yes, but that was the case for how my implementation went.

6.4 Final Curves

My network seems to be performing as I would expect a normal large neural network to perform. There is overfitting, yes, but the loss, both validation and training, are going down at a lessening rate each epoch. Later epochs had some odd overfitting on both datasets (CIFAR/MNIST).

6.5 Final Notes

This series was an extremely fun, but also sometimes frustating series. The abstraction of information was made easier to debug with the unit tests, but the original design was tedious and annoying to debug as well. The backprop is not necessarily stable, but with the right starting values, results in the current training that seems to perform well. Overall, one of my favorite lab series I have run.